Post-Generation NLI Entailment and Citation Verification.

Honovich et al. (2022) established the TRUE benchmark across 11 datasets and proved that binary Natural Language Inference entailment scoring outperforms ROUGE and BLEU at detecting factual inconsistencies. Min et al. (2023) introduced FActScore to decompose long-form generations into atomic claims and compute the exact proportion of supported facts.

Chern et al. (2023) demonstrated via FacTool that language models frequently fabricate nonexistent paper titles, author surnames, and wrong publication years during scientific literature synthesis. Qian et al. (2025) presented VeriCite, a three-stage framework that verifies each statement against its cited passage via NLI and prunes unentailed claims before synthesis.